

Akasa Air - Data Engineering Assignment

This is customer and order analytics, developed using the Medallion Architecture. The platform ingests data in multiple formats-CSV and XML, performs cleaning over time, and disseminates useful business KPIs through Gold-layer dashboards. Solutions are provided in both in-memory (Pandas) and database (MySQL) environments for more versatility.


- **Table-Based (MySQL) Approach:** SQL-driven processing to ensure scalability and high performance.
- **In-Memory (Pandas) Approach:** Makes use of Python for lightweight analytics, enabling rapid prototyping and ease of integration.

Data Sources

Source	Format	Fields
Customer Data	CSV	customer_id, customer_name, mobile_number, region
Order Data	XML	order_id, mobile_number, order_date_time, sku_id, sku_count, total_amount

Technical Architecture

Medallion Architecture Implementation



Project Structure

```
akasa-data-engineer/
├── main.py                      # Main pipeline execution script
├── requirements.txt               # Python dependencies
├── .env.example                  # Environment variables template
└── .gitignore                     # Git ignore rules

└── assets/                        # Static assets and outputs
    ├── task_DE_new_customers.csv   # Sample customer data
    ├── task_DE_new_orders.xml     # Sample order data
    └── analytics_dashboards/
        ├── dashboard_20251107_073641.png
        ├── revenue_20251107_073641.png
        └── regional_20251107_073641.png
    └── reports/                   # Business reports
        ├── business_report_20251107_043727.txt
        └── business_report_20251107_043727.json

└── data/                          # Processed data storage
    ├── bronze/                    # Raw ingested data
        ├── customers_bronze_20251106_230722.parquet
        └── orders_bronze_20251106_230722.parquet
    ├── silver/                    # Cleaned and enriched data
        ├── customers_silver_20251106_230722.parquet
        ├── orders_silver_20251106_230722.parquet
        └── processing_metadata_20251106_230722.json
    └── gold/                      # Business KPIs and metrics
        ├── repeat_customers_20251106_230722.parquet
        ├── monthly_trends_20251106_230722.parquet
        ├── regional_revenue_20251106_230722.parquet
        └── top_customers_30d_20251106_230722.parquet
        └── gold_metadata_20251106_230722.json

└── logs/                         # Application logs
    └── akasaair_processing.log      # Main processing log file

└── src/                          # Source code
    ├── __init__.py
    ├── config/
        ├── __init__.py
        └── config_manager.py       # Paths, logging, environment config
    ├── bronze/                    # Bronze Layer - Raw Data Ingestion
        ├── __init__.py
        ├── data_loader.py          # Main bronze data loader
        ├── csv_ingestor.py         # CSV customer data ingestion
        └── xml_ingestor.py         # XML order data ingestion
    ├── silver/                    # Silver Layer - Data Cleaning & Enrichment
        ├── __init__.py
        ├── silver_processor.py    # Main silver processor
        ├── data_cleaner.py         # Data cleaning operations
        ├── data_validator.py       # Data validation & quality checks
        └── data_enricher.py         # Feature engineering & enrichment
    ├── gold/                      # Gold Layer - Business Intelligence
        ├── __init__.py
        ├── gold_processor.py      # Main gold processor
        ├── kpi_calculator.py       # Core KPI calculations
        └── business_metrics.py     # Additional business metrics
    ├── presentation/              # Visualization & Reporting
        ├── __init__.py
        ├── visualizer.py          # Business dashboard generation
        └── report_generator.py      # Text and JSON report generation
    └── utils/                     # Utility functions
        ├── __init__.py
        ├── database.py            # Database connection & operations
        └── helpers.py              # Common helper functions

└── docs/                         # Documentation
```

Setup and Installation

Prerequisites

- Python 3.8+ (core data pipeline engine).
- MySQL 8.0+ (if choosing database approach).
- Dependency management via `pip install -r requirements.txt`.

Steps

1. Clone the project repository.
2. Install all dependencies.
3. If using database, configure `.env` with DB credentials.
4. Place source CSV/XML files in `./assets` directory.

Run `python main.py` for pipeline execution

- The script:
 - Loads customer (CSV) and order (XML) data.
 - Runs the Table-Based Approach (if MySQL credentials provided).
 - Runs the In-Memory Approach (Pandas).
 - Displays KPI results and logs the workflow to `data_processing.log`.

Sample Output

```

REPEAT CUSTOMERS:
-----
customer_name  number_of_orders
Aarav Mehta      2

MONTHLY TRENDS:
-----
month  total_orders  total_revenue
2025-09          1        35720
2025-10          1        22350
2025-11          1        15897

REGIONAL REVENUE:
-----
region  regional_revenue
West      38247
North    35720

TOP CUSTOMERS 30D:
-----
customer_name  recent_spend
Aarav Mehta      38247

VISUALIZATIONS GENERATED:
-----
• Comprehensive Dashboard: assets\analytics_dashboards\dashboard_20251107_073951.png
• Revenue: assets\analytics_dashboards\revenue_20251107_073951.png
• Regional: assets\analytics_dashboards\regional_20251107_073951.png

=====
KPI CALCULATION & VISUALIZATION COMPLETE
=====
INFO: Gold layer completed: All business KPIs calculated
INFO: =====
INFO: TABLE-BASED APPROACH: SQL Database Processing
INFO: =====
INFO: All KPIs calculated securely using parameterized SQL queries

```

Core Components

Dual Processing Engine

Approach	Use Case	Advantages
In-Memory (Pandas)	Rapid development & testing	Zero dependencies, instant execution
Database (MySQL)	Production workloads	Scalability, ACID compliance, SQL power

Bronze Layer - Raw Data Ingestion

- **CSVIngestor:** Customer data processing with schema validation
- **XMLIngestor:** Order data extraction with structure preservation
- **DataLoader:** Coordination of multi-source data ingestion

Silver Layer - Data Refinement

- **DataCleaner:** Data quality enforcement & standardization
- **DataValidator:** Comprehensive validation rules & integrity checks
- **DataEnricher:** Feature engineering & business context addition

Gold Layer - Business Intelligence

- **KPICalculator:** Core KPI computation engine
- **BusinessMetrics:** Advanced analytics & customer segmentation
- **Visualization Engine:** Professional dashboard generation
- **ReportGenerator:** Multi-format business reporting

Security and Best Practices

- All SQL queries use parameterization through SQLAlchemy, protecting against injection issues.
- Credentials are never exposed in code or logs.
- All input data is validated and sanitized before processing.
- Operational events are logged in data_processing.log for auditability.

Conclusion

This project effectively delivers a strong, dual-approach data processing pipeline: database and in-memory computation for creating business-critical KPIs. The architecture promotes modularity, scalability, and strict data security.